


Meta-learning for Reinforcement Learning

Since REINFORCE is based on gradient-based optimization of $J(\theta)$, both MAML-like and Model-based Meta-learning methods apply *directly*, as also demonstrated in those papers.

Reference - [Learning to Reinforcement Learn](#) and - [RL² Fast Reinforcement Learning via Slow Reinforcement Learning](#)

RL tasks are sampled from a distribution over MDPs, and are presented as sequences, as in model-based meta-learning. Instead of y_{i-1} , here a_{i-1}, r_{i-1} are passed along with s_i . First paper uses advantage AAC and second uses TRPO, to train an RNN as the deep-RL network; otherwise similar. Hidden state is re-set every time a new episode starts (first paper) / when a new MDP is sampled (second paper, better).



Meta-learning for Continual Learning

Let A_i^k = training accuracy on task k after encountering i tasks, and T_i^k the ‘time’ to ‘learn’ task k . Then $CF = \sum_{i>N} \sum_{k<N} (A_k^k - A_i^k)$ is a measure of “forgetting”; and $FT = \sum_{t<N} \sum_{s>N} (T_t^t - T_s^s)$ measures “forward transfer”.

In a meta-learning setup for continual learning, each task has a train and test set. The first N tasks are considered meta-training and the rest meta-testing. During meta-training, all tasks are available; in meta-testing, only the most recent task is available. If \hat{T}_i^k the time to learn on the test-set for task k after seeing i tasks, then $CT_i = \sum_{k<n} (\hat{T}_i^k - \hat{T}_k^k)$ for $i > N$ also measures ‘forgetting’. Further, if \hat{A}_k^k is the **test**-set accuracy $G = \sum_{k>N} \hat{A}_N^k$ for $k > N$ measures generalization from meta-learning. Often, continual learning *trajectories* CT_i , $CF_i = \sum_{k<N} (A_k^k - A_i^k)$, $FT_i = \sum_{t<N} (T_t^t - T_i^i)$ and \hat{A}_N^i , for $i > N$ are measured.

Continual Learning: Definition, & Gated Linear Networks

Continual learning: (a) avoid “forgetting” previously learned tasks, and also (b) improve as new tasks are learned. Online: single pass over data/tasks.

References - [Gated Linear Networks](#) and [A Combinatorial Perspective on Transfer Learning](#)

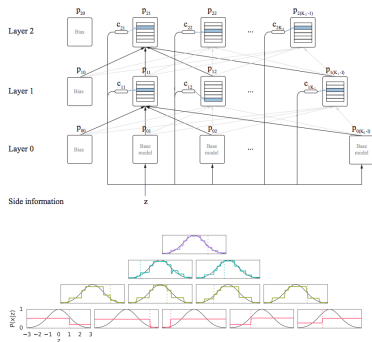
GLNs: Each neuron predicts the target directly, i.e., no back-propagation, via a **geometric mixture** of probabilities obtained from the previous layer. Lowest layer can be *random*, with say K_0 neurons. Each neuron in layer i has 2^d weight vectors w_{ikc} (of size K_{i-1}). Which weight vector to use is determined by a hash $c_{ik}(z)$ ('half-space gating') of the input z .

Outputs are: $p_{ik}(z) = \sigma(w_{ikc_{ik}(z)}^T \cdot \sigma^{-1}(p_{i-1}(z)))$

& updated as: $\Delta w_{ikc_{ik}(z)} = \eta(p_{ik} - y)\sigma^{-1}(p_{i-1})$


For combinatorial generalization, a Bayesian mixture of previous layers' models from past tasks is chosen at each neuron.

Note: $\sigma(x) = \frac{1}{1+e^{-x}}$ and $\sigma^{-1}(x) = \log \frac{x}{1-x}$




Continual Learning: Supermasks in Superposition

References - [Supermasks in Superposition](#) and [What's Hidden in a Randomly Weighted Neural Network?](#)



During training learn a mask M per task, using top $k\%$ scores as above; network outputs $p = f(x, W \odot M)$.



For future tasks adapt a mixture of masks by gradient descent on **entropy**: i.e., update $p(\alpha) = f(W \odot \sum_i \alpha_i M^i)$ via $\alpha \leftarrow \alpha - \eta \nabla_{\alpha} \mathcal{H}(p(\alpha))$, so low entropy, i.e., less 'confused' masks are preferred.